

Boosting Linux Storage Performance with LVM Striping



Ahmed Mansouri · Follow

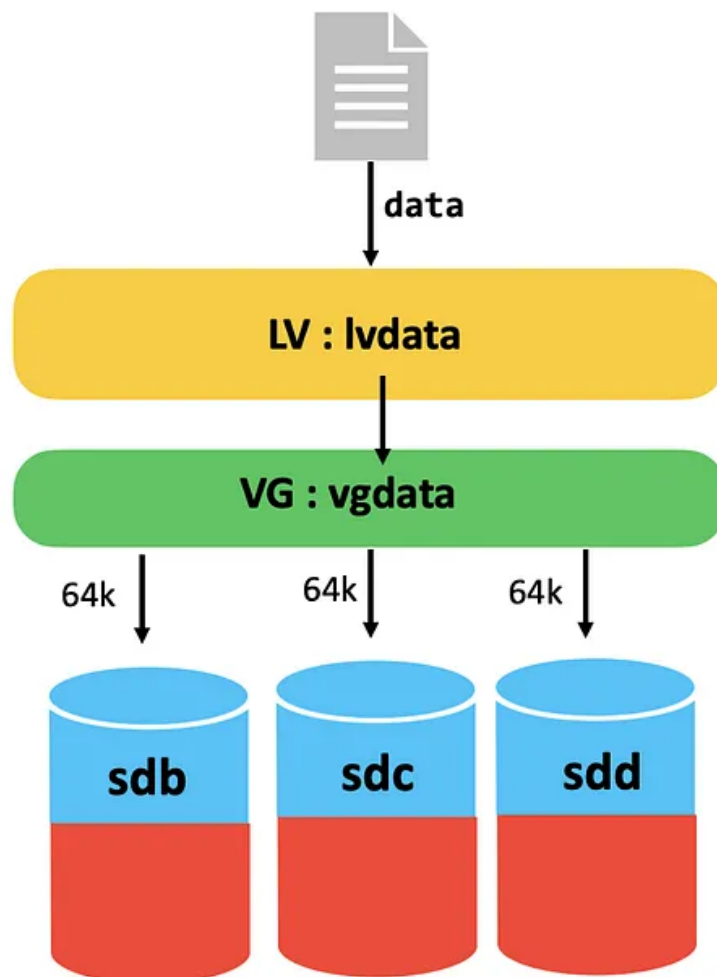
6 min read · Feb 4, 2024



Listen



Share




In the realm of storage management, achieving optimal disk performance takes precedence. The need for faster data access, reduced latency, and improved I/O operations and throughput has led to the adoption of advanced techniques. One such technique is LVM Striping, a powerful feature of the Logical Volume Manager LVM that can significantly enhance volume performance.

I — Linear LVM vs Striping LVM:

I.1 — Linear LVM:


The linear configuration, often considered the **standard LVM** setup, involves adding multiple physical volumes (disks) to a volume group in a linear fashion. This means that data is sequentially stored across these volumes, utilizing one disk before moving on to the next.



While linear LVM provides a straightforward approach to storage expansion, it may not fully exploit the potential for parallel processing and increased throughput.

I.2 — Striping LVM:

In contrast, Striping LVM takes a more advanced approach by distributing data across multiple physical volumes simultaneously. This striping process creates a logical volume that spans the disks, allowing for parallel read and write operations. The result is a substantial boost in performance, making Striping LVM an attractive option for environments with high I/O and throughput demands.



1. Stripe Creation:

- The data is divided into segments known as “stripes”
- Each stripe is written to a specific disk in the striped logical volume LV.

2. Parallel Writing of Stripes:

- These stripes are simultaneously written across the multiple physical volumes PVs (disks).

3. Balanced Distribution:

- The stripes are distributed evenly across the disks, preventing a single disk from becoming a bottleneck.

4. Enhanced IOPS and Throughput:

- Through the parallel writing of stripes, LVM Striping significantly boosts the overall throughput and performance of the logical volume.

I.3 — Scenario :

Consider a scenario with three disk drives allocated to three physical volumes PVs. If each individual physical volume can achieve a total of 125M/s as throughput :

- Employing “**LVM Striping**” would result in a volume group capable of 375M/s .
- In contrast, using “**LVM Linear**”, the throughout remains at 125M/s , and no matter how many disks we add in LVM.

II — Setting up “Linear LVM” and “Striping LVM”:

In this section, we’ll establish two volume groups (VGs): the initial one for Linear LVM utilizing 3 disks, and the second one for Striping LVM, also with three disks. Following the VG setup, we’ll proceed to create Logical Volumes (LVs) on each group, then formatting and mounting them

II.1 — Initial check

Environment :

- Server : AWS EC2 instance = *m4.10xlarge* ,
- EBS Volumes : *6 x 20G EBS volumes* .
- Each EBS volume has the following characteristics : { Type : *GP3*, IOPS = *3000* , Throughput = *125MiB/s* }
- OS/image : Amazon Linux 2

Here the list of the volumes :

											< 1 > ⚙	
<input type="checkbox"/>	Name ▾	Volume ID ▾	Type ▾	Size ▾	IOPS ▾	Throughput ▾	Snapshot ▾	Created				
<input type="checkbox"/>	-	vol- [REDACTED] 15	gp3	20 GiB	3000	125	-	2024/02/03 23:07				
<input type="checkbox"/>	-	vol- [REDACTED] 7	gp3	20 GiB	3000	125	-	2024/02/03 23:07				
<input type="checkbox"/>	-	vol- [REDACTED] 4c	gp3	20 GiB	3000	125	-	2024/02/03 23:07				
<input type="checkbox"/>	-	vol- [REDACTED] b	gp3	20 GiB	3000	125	-	2024/02/03 23:07				
<input type="checkbox"/>	-	vol- [REDACTED] b	gp3	20 GiB	3000	125	-	2024/02/03 23:07				
<input type="checkbox"/>	-	vol- [REDACTED] 77	gp3	20 GiB	3000	125	-	2024/02/03 23:07				

Let’s check our EC2 instance and see our disks

```
# lsblk
```

```
[root@ip-172-31-4-125 ec2-user]# lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT
xvda        202:0    0   8G  0 disk
└─xvda1     202:1    0   8G  0 part /
xvdb        202:16   0  20G  0 disk
xvdc        202:32   0  20G  0 disk
xvdd        202:48   0  20G  0 disk
xvde        202:64   0  20G  0 disk
xvdf        202:80   0  20G  0 disk
xvdg        202:96   0  20G  0 disk
[root@ip-172-31-4-125 ec2-user]#
```

II.2 — Setting up the VGs:

In order to be able to create the LVs, we need first to create the VG for each one of them :

```
# vgcreate vg_linear /dev/sdb /dev/sdc /dev/sdd

# vgcreate vg_stripping /dev/sde /dev/sdf /dev/sdg
```

```
[root@ip-172-31-4-125 ec2-user]# vgcreate vg_linear /dev/sdb /dev/sdc /dev/sdd
Physical volume "/dev/sdb" successfully created.
Physical volume "/dev/sdc" successfully created.
Physical volume "/dev/sdd" successfully created.
Volume group "vg_linear" successfully created
[root@ip-172-31-4-125 ec2-user]#
[root@ip-172-31-4-125 ec2-user]# vgcreate vg_stripping /dev/sde /dev/sdf /dev/sdg
Physical volume "/dev/sde" successfully created.
Physical volume "/dev/sdf" successfully created.
Physical volume "/dev/sdg" successfully created.
Volume group "vg_stripping" successfully created
[root@ip-172-31-4-125 ec2-user]#
```

```
# vgs
```

```
[root@ip-172-31-4-125 ec2-user]# vgs
VG          #PV #LV #SN Attr   VSize  VFree
vg_linear   3   0   0 wz--n-  <59.99g <59.99g
vg_stripping 3   0   0 wz--n-  <59.99g <59.99g
[root@ip-172-31-4-125 ec2-user]#
```

II.3 — Create a Logical Volume as “Linear LVM”

```
# lvcreate -l 100%FREE -n lv_linear vg_linear
```

```
[root@ip-172-31-4-125 ec2-user]#
[root@ip-172-31-4-125 ec2-user]# lvcreate -l 100%FREE -n lv_linear vg_linear
Logical volume "lv_linear" created.
[root@ip-172-31-4-125 ec2-user]#
[root@ip-172-31-4-125 ec2-user]#
```

- `-l 100%FREE` : The size of LV is set to match the entirety of the available free space within the VG `vg_linear` .
- `-n lv_linear` : LV name
- `vg_linear` : Target VG

II.4 — Create a Logical Volume as “Striping LVM”

```
# lvcreate -l 100%FREE -i 3 -I 64k -n lv_stripping vg_stripping
```

```
[root@ip-172-31-4-125 ec2-user]#
[root@ip-172-31-4-125 ec2-user]# lvcreate -l 100%FREE -i 3 -I 64k -n lv_stripping vg_stripping
Logical volume "lv_stripping" created.
[root@ip-172-31-4-125 ec2-user]#
```

- `-l 100%FREE` : The size of LV is set to match the entirety of the available free space within the VG `vg_stripping` .
- `-n lv_stripping` : LV name
- `vg_stripping` : Target VG
- `-i 3` : stripes Number
- `-I 64k` : stripe size

II.5 — Check the LVs

Let's check the new LVs:

- Let's start by `lv_linear`

```
# lvdisplay -m /dev/vg_linear/lv_linear
```

```
[root@ip-172-31-4-125 ec2-user]#  
[root@ip-172-31-4-125 ec2-user]# lvdisplay -m /dev/vg_linear/lv_linear  
--- Logical volume ---  
LV Path                /dev/vg_linear/lv_linear  
LV Name                 lv_linear  
VG Name                 vg_linear  
LV UUID                 r7b5QG-lsya-Pgrq-VY7M-SQsY-Fo8r-PRF0kB  
LV Write Access         read/write  
LV Creation host, time ip-172-31-4-125.eu-west-1.compute.internal, 2024-02-03 20:21:15 +0000  
LV Status                available  
# open                  0  
LV Size                  <59.99 GiB  
Current LE               15357  
Segments                3  
Allocation               inherit  
Read ahead sectors      auto  
- currently set to      256  
Block device             253:0  
  
--- Segments ---  
Logical extents 0 to 5118:  
  Type                   linear  
  Physical volume         /dev/sdb  
  Physical extents        0 to 5118  
  
Logical extents 5119 to 10237:  
  Type                   linear  
  Physical volume         /dev/sdc  
  Physical extents        0 to 5118  
  
Logical extents 10238 to 15356:  
  Type                   linear  
  Physical volume         /dev/sdd  
  Physical extents        0 to 5118
```

- Let's check new `lv_striping`

```
# lvdisplay -m /dev/vg_striping/lv_striping
```

```
[root@ip-172-31-4-125 ec2-user]# lvdisplay -m /dev/vg_stripping/lv_stripping
--- Logical volume ---
LV Path                /dev/vg_stripping/lv_stripping
LV Name                 lv_stripping
VG Name                 vg_stripping
LV UUID                ESuJWf-262j-wVG6-5zmC-NX0V-5YDD-Gkz74R
LV Write Access         read/write
LV Creation host, time ip-172-31-4-125.eu-west-1.compute.internal, 2024-02-03 20:28:34 +0000
LV Status                available
# open                  0
LV Size                 <59.99 GiB
Current LE              15357
Segments                1
Allocation              inherit
Read ahead sectors      auto
  - currently set to    768
Block device            253:1

--- Segments ---
Logical extents 0 to 15356:
Type                striped ←
Stripes              3
Stripe size          64.00 KiB
Stripe 0:
  Physical volume     /dev/sde
  Physical extents    0 to 5118
Stripe 1:
  Physical volume     /dev/sdf
  Physical extents    0 to 5118
Stripe 2:
  Physical volume     /dev/sdg
  Physical extents    0 to 5118
```

==> Here we can see the different details about the striping that we configured .

II.6 — Formating and mounting the LVs

```
# mkfs.xfs /dev/vg_linear/lv_linear

# mkfs.xfs /dev/vg_stripping/lv_stripping
```

```
# mkdir /mnt/linear

# mkdir /mnt/stripping
```

```
# mount /dev/vg_linear/lv_linear /mnt/linear/

# mount /dev/vg_stripping/lv_stripping /mnt/stripping/
```



```
# df -h
```

```
[root@ip-172-31-4-125 ec2-user]#  
[root@ip-172-31-4-125 ec2-user]# df -h  
Filesystem      Size  Used Avail Use% Mounted on  
devtmpfs        1.9G   0    1.9G   0% /dev  
tmpfs           2.0G   0    2.0G   0% /dev/shm  
tmpfs           2.0G 448K    2.0G   1% /run  
tmpfs           2.0G   0    2.0G   0% /sys/fs/cgroup  
/dev/xvda1      8.0G  1.7G   6.4G  21% /  
tmpfs           391M   0    391M   0% /run/user/1000  
/dev/mapper/vg_striping-lv_striping 60G 462M   60G   1% /mnt/striping  
/dev/mapper/vg_linear-lv_linear     60G 461M   60G   1% /mnt/linear
```

==> The LVs are now mounted correctly under the appropriate directories

III — Benchmarks the LVs/disks

In order to benchmark our LVs/disks, we will use the `fio` tool. So let us install it before :

```
# yum install fio -y
```

III.1 — Benchmark “lv_linear”

In order to benchmark the LV, we will use the `fio` config file below that will help us generating traffic for 400M as throughput :

```
# cat fio_config-1.fio  
  
[global]  
ioengine=libaio  
runtime=60  
time_based  
direct=1  
rw=write  
size=10G  
bs=512K  
rate=400M  
numjobs=16
```

```
[job1]
filename=/mnt/linear/testfile
```

Execute now the `fio` tool as below using the defined file above:

```
# fio fio_config-1.fio
```

- In the same time, run the `iostat` tool to monitor the disks usage in a separate terminal :

```
# iostat -xdmt 2
```

```
[root@ip-172-31-4-125 ec2-user]# iostat -xdmt 2
Linux 5.10.205-195.807.amzn2.x86_64 (ip-172-31-4-125.eu-west-1.compute.internal)    03/02/24    _x86_64_    (2 CPU)

03/02/24 23:21:06
Device:            rrqm/s    wrqm/s      r/s      w/s    rMB/s    wMB/s avgrq-sz avgqu-sz   await  r_await  w_await   svctm   %util
xvda                0.00      0.00      0.00      0.00      0.00      0.00    0.00      0.00      0.00    0.00    0.00    0.00    0.00    0.00
xvdb                0.00      0.00      0.00     501.00      0.00    125.25   512.00    31.11    62.10    0.00    62.10    2.00   100.00
xvdc                0.00      0.00      0.00      0.50      0.00      0.00     5.00      0.00      0.00    0.00    0.00    4.00     0.20
xvdd                0.00      0.00      0.00      0.00      0.00      0.00     0.00      0.00      0.00    0.00    0.00    0.00     0.00
xvde                0.00      0.00      0.00      0.00      0.00      0.00     0.00      0.00      0.00    0.00    0.00    0.00     0.00
xvdf                0.00      0.00      0.00      0.00      0.00      0.00     0.00      0.00      0.00    0.00    0.00    0.00     0.00
xvdg                0.00      0.00      0.00      0.00      0.00      0.00     0.00      0.00      0.00    0.00    0.00    0.00     0.00
dm-0                0.00      0.00      0.00    251.00      0.00    125.25  1021.97    15.99    63.72    0.00    63.72    3.98   100.00
dm-1                0.00      0.00      0.00      0.00      0.00      0.00     0.00      0.00      0.00    0.00    0.00    0.00     0.00

03/02/24 23:21:08
Device:            rrqm/s    wrqm/s      r/s      w/s    rMB/s    wMB/s avgrq-sz avgqu-sz   await  r_await  w_await   svctm   %util
xvda                0.00      0.00      0.00      0.00      0.00      0.00     0.00      0.00      0.00    0.00    0.00    0.00     0.00
xvdb                0.00      0.00      0.00     500.50      0.00    125.12   512.00    31.09    62.12    0.00    62.12    2.00   100.20
xvdc                0.00      0.00      0.00      0.00      0.00      0.00     0.00      0.00      0.00    0.00    0.00    0.00     0.00
xvdd                0.00      0.00      0.00      0.00      0.00      0.00     0.00      0.00      0.00    0.00    0.00    0.00     0.00
xvde                0.00      0.00      0.00      0.00      0.00      0.00     0.00      0.00      0.00    0.00    0.00    0.00     0.00
xvdf                0.00      0.00      0.00      0.00      0.00      0.00     0.00      0.00      0.00    0.00    0.00    0.00     0.00
xvdg                0.00      0.00      0.00      0.00      0.00      0.00     0.00      0.00      0.00    0.00    0.00    0.00     0.00
dm-0                0.00      0.00      0.00    250.50      0.00    125.25  1024.00    16.01    63.90    0.00    63.90    4.00   100.20
dm-1                0.00      0.00      0.00      0.00      0.00      0.00     0.00      0.00      0.00    0.00    0.00    0.00     0.00
```

==> We can see from the output above that ONLY the first disk `xvdb` is used and that the throughput is equal to 125M which is the baseline value of one EBS volume itself.

The same thing is showing by the summary given the `fio` tool :

```
Run status group 0 (all jobs):
  WRITE: io=7641.0MB, aggrb=130267KB/s, minb=8141KB/s, maxb=8152KB/s, mint=60003msec, maxt=60064msec

Disk stats (read/write):
dm-0: ios=0/15267, merge=0/0, ticks=0/957160, in_queue=957160, util=99.85%, aggrios=0/10190, aggrmerge=0/1, aggrticks=0/621926, aggrin_queue=621927, aggrutil=99.68%
xvda: ios=0/3, merge=0/0, ticks=0/2, in_queue=2, util=0.02%
xvdd: ios=0/0, merge=0/0, ticks=0/0, in_queue=0, util=0.00%
xvdb: ios=0/30567, merge=0/4, ticks=0/1865778, in_queue=1865779, util=99.68%
```

==> WRITE operations were exclusively performed on a single disk.

III.2 — Benchmark “lv_striping”

Use the same `fio` config file above but just change the `filename` parameter :

```
filename=/mnt/striping/testfile
```

Use `iostat` to monitor the disks

```
[root@ip-172-31-4-125 ec2-user]# iostat -xmt 2
```

Linux 5.10.205-195.807.amzn2.x86_64 (ip-172-31-4-125.eu-west-1.compute.internal) 03/02/24 _x86_64_ (2 CPU)

03/02/24 23:18:55

Device:	rrqm/s	wrqm/s	r/s	w/s	rMB/s	wMB/s	avgrq-sz	avgqu-sz	await	r_await	w_await	svctm	%util
xvda	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
xvdb	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
xvdc	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
xvdd	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
xvde	0.00	1250.00	0.00	750.50	0.00	125.03	341.19	15.93	21.23	0.00	21.23	1.33	99.80
xvdf	0.00	1252.50	0.00	751.00	0.00	125.28	341.65	2.89	3.85	0.00	3.85	1.33	100.00
xvdg	0.00	1250.00	0.00	751.00	0.00	125.12	341.22	1.42	1.89	0.00	1.89	1.33	100.00
dm-0	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
dm-1	0.00	0.00	0.00	6004.00	0.00	375.25	128.00	54.85	9.14	0.00	9.14	0.17	100.00

03/02/24 23:18:57

Device:	rrqm/s	wrqm/s	r/s	w/s	rMB/s	wMB/s	avgrq-sz	avgqu-sz	await	r_await	w_await	svctm	%util
xvda	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
xvdb	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
xvdc	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
xvdd	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
xvde	0.00	1255.50	0.00	752.50	0.00	125.34	341.13	15.96	21.21	0.00	21.21	1.33	100.20
xvdf	0.00	1254.00	0.00	750.50	0.00	125.09	341.36	3.07	4.09	0.00	4.09	1.33	99.80
xvdg	0.00	1253.00	0.00	752.00	0.00	125.25	341.11	1.47	1.96	0.00	1.96	1.33	100.00
dm-0	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
dm-1	0.00	0.00	0.00	6020.00	0.00	376.25	128.00	55.76	9.26	0.00	9.26	0.17	100.00

==> All 3 disks used by the `lv_striping` participated in the WRITE operations, resulting in a combined throughput equal to the sum of the individual disk throughputs ($125\text{M} \times 3 = 375\text{M}$).

The same thing is showing by the summary given the `fio` tool :

```
Run status group 0 (all jobs):
  WRITE: io=21535MB, aggrbw=367426KB/s, minb=22951KB/s, maxb=22990KB/s, mint=60002msec, maxt=60017msec

Disk stats (read/write):
  dm-1: ios=0/344040, merge=0/0, ticks=0/3262664, in_queue=3262664, util=99.82%, aggrios=0/43074, aggrmerge=0/71784, aggrticks=0/405246, aggrin_queue=405246, aggrutil=99.68%
  xvdf: ios=0/43074, merge=0/71784, ticks=0/177379, in_queue=177379, util=99.56%
  xvdg: ios=0/43073, merge=0/71782, ticks=0/82033, in_queue=82032, util=99.48%
  xvde: ios=0/43075, merge=0/71788, ticks=0/956328, in_queue=956328, util=99.68%
```

Conclusion:

LVM Striping stands as a robust solution for organizations seeking to unlock the full potential of their storage infrastructure. By distributing data intelligently across multiple disks, LVM Striping not only boosts performance but also provides a